

# Feature Extraction-Copy2

September 17, 2018

## 1 Helper Functions and Training Procedure:

```
In [2]: import cv2
import numpy as np
import matplotlib.pyplot as plt
from mnist import MNIST
from skimage.feature import hog
import itertools

In [13]: from sklearn import svm

In [17]: from sklearn import neighbors

In [3]: def generate_hog(X, cell_size=(8,8)):
    hog_img=[]
    for img in X:
        vec=hog(img, pixels_per_cell=cell_size)
        hog_img.append(vec)
    return np.asarray(hog_img)

In [4]: def Sigmoid(x):
    return np.asarray(1/(1+np.exp(-x)))

def ReLU(x):
    return np.asarray(np.maximum(0,x))

def Softmax(x):
    b=np.exp(x)
    c=np.sum(b, axis=0)
    d=np.divide(b, c)
    return d

def ReLU_grad(x):
    return np.heaviside(x,0)

def Sigmoid_grad(x):
    return np.asarray(np.multiply(Sigmoid(x),1-Sigmoid(x)))
```

```

def CrossEntropy(target_, output_):
    return np.sum(-np.sum(np.multiply(target_, np.log((output_))), axis=0))

def CrossEntropy_grad(target_, output_):
    return -np.divide(target_, output_)

def Softmax_CE_grad(target_, pred_):
    grad=-pred_+target_
    return grad

In [5]: def add_noise(images):
        size=images.shape
        x=np.random.normal(loc=0.0, scale=10, size=size)
        noisy=np.add(images,x)
        return noisy

def labels_to_class(labels):
    return np.argmax(labels, axis=1)

def confusion_matrix(target_, pred):
    size=len(target_[0])
    target_class=labels_to_class(target_)
    pred_class=labels_to_class(pred)
    cm=np.zeros([size, size])
    for a,p in zip(target_class, pred_class):
        cm[a][p]+=1
    return cm

def cm_metrics(cm):
    diag=(np.diagonal(cm))
    psum=np.sum(cm, axis=0, dtype=np.float32)
    rsum=np.sum(cm, axis=1, dtype=np.float32)
    p=np.divide(diag, psum)
    r=np.divide(diag, rsum)
    prod=np.multiply(p,r)
    sum_=p+r
    f=2*np.divide(prod,sum_)
    a=1.*np.sum(diag)/np.sum(cm)
    return (a,p,r,f)

def plot_confusion_matrix(cm, target_names, title='Confusion matrix',):

    cmap = plt.get_cmap('Blues')

    plt.figure(figsize=(8, 6))
    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)

```

```

plt.colorbar()

tick_marks = np.arange(len(target_names))
plt.xticks(tick_marks, target_names)
plt.yticks(tick_marks, target_names)

thresh = cm.max() / 2
for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
    plt.text(j, i, "{:,}".format(cm[i, j]),
             horizontalalignment="center",
             color="white" if cm[i, j] > thresh else "black")

plt.tight_layout()
plt.ylabel('True label')
plt.xlabel('Predicted label')
plt.show()

def plot_loss(train_loss, test_loss, title):
    x=200*np.arange(0, len(train_loss))
    plt.plot(x,train_loss, label='Train')
    plt.plot(x,test_loss, label='Test')
    plt.title(title)
    plt.xlabel("Iterations")
    plt.ylabel("Average Loss")
    plt.legend(loc='upper left')
    plt.show()

def top_k_pred(pred, k):
    sort=np.argsort(pred, axis=0)
    return sort[::-1][:k]

In [6]: def shuffle_data(images, labels):
        index=np.random.permutation(len(images))
        shuff_images, shuff_labels=images[index], labels[index]
        return shuff_images, shuff_labels

def get_split_mask(images):
    size=len(images)/5
    a=np.ones(5*size, dtype=bool)
    b=np.arange(10)
    mask=[]
    for i in range(5):
        mask.append([False if (j<(i+1)*size)&(j>=(i)*size) else x for j,x in enumerate(a)
    return mask

def get_split_data(mask, images, labels, index):
    mask_=mask[index]
    train_images=images[mask_]

```

```

train_labels=labels[mask_]
inv_mask_=np.invert(mask_)
test_images=images[inv_mask_]
test_labels=labels[inv_mask_]
return train_images, train_labels, test_images, test_labels

def get_accuracy(model, test_images, test_labels):
    batch_size=len(test_images[0])
    count=np.zeros(batch_size)
    pred=model.forward(test_images)
    count=[1 if np.argmax(pred[:,i], axis=0)==np.argmax(test_labels[:,i], axis=0) else 0
    correct=np.sum(count)
    accuracy=100*correct/batch_size
    return(pred, accuracy)

def SGD_mom(model, batch_images, batch_labels, l2):
    batch_size=len(batch_images[0])
    loss=model.update(batch_images, batch_labels, l2)
    return loss

def one_hot(labels):
    a=np.zeros((len(labels), 10))
    a[np.arange(len(labels)),labels]=1
    return a

```

```

In [7]: class MLP(object):
    def __init__(self, input_size, h1_size, h2_size, output_size):
        self.W1=np.random.normal(loc=0.0, scale=0.08, size=(h1_size, input_size) )
        self.W2=np.random.normal(loc=0.0, scale=0.08, size=(h2_size, h1_size) )
        self.W3=np.random.normal(loc=0.0, scale=0.08, size=(output_size, h2_size) )
        self.B1=np.zeros(h1_size).reshape(-1,1)
        self.B2=np.zeros(h2_size).reshape(-1,1)
        self.B3=np.zeros(output_size).reshape(-1,1)

        self.W1_grad=np.zeros_like(self.W1)
        self.W2_grad=np.zeros_like(self.W2)
        self.W3_grad=np.zeros_like(self.W3)
        self.B1_grad=np.zeros_like(self.B1)
        self.B2_grad=np.zeros_like(self.B2)
        self.B3_grad=np.zeros_like(self.B3)

        self.W1_mom=np.zeros_like(self.W1)
        self.W2_mom=np.zeros_like(self.W2)
        self.W3_mom=np.zeros_like(self.W3)
        self.B1_mom=np.zeros_like(self.B1)
        self.B2_mom=np.zeros_like(self.B2)
        self.B3_mom=np.zeros_like(self.B3)

```

```

def update(self, input_, target_, l2=0):

    batch_size=len(input_[0])
    x=input_
    h1=np.add(np.matmul(self.W1, input_), self.B1)
    a1=act[act_ind](h1)
    h2=np.add(np.matmul(self.W2, a1), self.B2)
    a2=act[act_ind](h2)
    h3=np.add(np.matmul(self.W3, a2), self.B3)
    a3=Softmax(h3)

    loss=CrossEntropy(target_, a3)
    _E_h3=Softmax_CE_grad(a3, target_)

    _a2_h2=act_grad[act_ind](h2)
    _a1_h1=act_grad[act_ind](h1)

    _E_W3=np.matmul(_E_h3,np.transpose(a2))
    _E_B3=np.sum(_E_h3, axis=1).reshape(-1,1)
    _E_a2=np.matmul(np.transpose(self.W3), _E_h3)

    _E_h2=np.multiply(_E_a2, _a2_h2)
    _E_W2=np.matmul(_E_h2, np.transpose(a1))
    _E_B2=np.sum(_E_h2, axis=1).reshape(-1,1)
    _E_a1=np.matmul(np.transpose(self.W2), _E_h2)

    _E_h1=np.multiply(_E_a1, _a1_h1)
    _E_W1=np.matmul(_E_h1, np.transpose(x))
    _E_B1=np.sum(_E_h1, axis=1).reshape(-1,1)
    _E_x=np.matmul(np.transpose(self.W1), _E_h1)

    self.W1_grad=_E_W1/batch_size+self.W1*2*l2
    self.W2_grad=_E_W2/batch_size+self.W2*2*l2
    self.W3_grad=_E_W3/batch_size+self.W3*2*l2
    self.B1_grad=_E_B1/batch_size
    self.B2_grad=_E_B2/batch_size
    self.B3_grad=_E_B3/batch_size

    self.W1_mom=gamma*self.W1_mom+lr*self.W1_grad
    self.W2_mom=gamma*self.W2_mom+lr*self.W2_grad
    self.W3_mom=gamma*self.W3_mom+lr*self.W3_grad
    self.B1_mom=gamma*self.B1_mom+lr*self.B1_grad
    self.B2_mom=gamma*self.B2_mom+lr*self.B2_grad
    self.B3_mom=gamma*self.B3_mom+lr*self.B3_grad

    self.W1-=self.W1_mom
    self.W2-=self.W2_mom

```

```

        self.W3-=self.W3_mom
        self.B1=self.B1-self.B1_mom
        self.B2=self.B2-self.B2_mom
        self.B3=self.B3-self.B3_mom

    return loss

    def forward(self, input_):
        x=input_
        h1=np.add(np.matmul(self.W1, input_), self.B1)
        a1=act[act_ind](h1)
        h2=np.add(np.matmul(self.W2, a1), self.B2)
        a2=act[act_ind](h2)
        h3=np.add(np.matmul(self.W3, a2), self.B3)
        a3=Softmax(h3)
        return a3

```

```

In [28]: input_size=81
        h1_size=500
        h2_size=250
        output_size=10
        gamma=0.99
        lr=1e-3
        act_ind=0
        batch_size=64
        epochs=10
        act=[Sigmoid, ReLU]
        act_grad=[Sigmoid_grad, ReLU_grad]
        actstr={0: 'Sigmoid', 1: 'ReLU'}

```

```

In [9]: from mnist import MNIST
        data=MNIST('/home/pradeep/data/')
        images, labels_ = data.load_training()

        images=np.asarray(images)
        labels_=np.asarray(labels_)
        images=images.reshape(-1,28,28)

        act=[Sigmoid, ReLU]
        act_grad=[Sigmoid_grad, ReLU_grad]

        images, labels_=shuffle_data(images, labels_)
        labels=one_hot(labels_)
        mask=get_split_mask(images)
        hog_img=generate_hog(X=images)
        hog_img*=255

```

/home/pradeep/.local/lib/python2.7/site-packages/skimage/feature/\_hog.py:150: skimage\_deprecatio

skimage\_deprecation)

```
In [27]: def train(model, epochs, images, labels, fold_index, l2=0):
    train_loss=[]
    test_loss=[]
    train_images, train_labels, test_images, test_labels=get_split_data(mask, images, l
    num_batches=len(train_images)/batch_size
    batch_images=np.array_split(train_images, num_batches)
    batch_labels=np.array_split(train_labels, num_batches)
    for epoch in range(epochs):
        for i in range(num_batches):
            size=len(batch_images[i])
            loss=SGD_mom(model, np.transpose(batch_images[i]), np.transpose(batch_label
            train_avg_loss=loss/size
            if((i)%200==0):
                print("Epoch "+str(epoch+1)+" Iteration "+str(i+1)+" : Avg Loss = "+str
                pred, accuracy=get_accuracy(model, np.transpose(test_images), np.transp
                test_avg_loss= CrossEntropy(np.transpose(test_labels), pred)/len(test_i
                train_loss.append(train_avg_loss)
                test_loss.append(test_avg_loss)
    plot_loss(train_loss, test_loss, "Plot of loss for "+actstr[act_ind]+" for Fold "+s
    output_cm_scores(model, test_images, test_labels)
    #plot_images(model, test_images[:20], 3)

In [10]: def output_cm_scores_(pred, labels):
    cm=confusion_matrix_(pred=pred, target_=labels)
    a,p,r,f=cm_metrics(cm)
    print "Accuracy = "+ str(a)
    print "Precision = ", p
    print "Recall = ", r
    print "F1 Score = ", f
    target_names=['0','1','2','3','4','5','6','7','8','9']
    plot_confusion_matrix(cm,target_names)

    def confusion_matrix_(target_, pred):
        size=10
        cm=np.zeros([size, size])
        for a,p in zip(target_, pred):
            cm[a][p]+=1
        return cm

In [18]: def train_KNN(hog_img, labels, fold_index):
    KNN=neighbors.KNeighborsClassifier(n_neighbors=5, n_jobs=-1)
    mask=get_split_mask(hog_img)
    train_hog, train_labels, test_hog, test_labels=get_split_data(images=hog_img, label
    KNN.fit(train_hog, train_labels)
    pred=KNN.predict(test_hog)
    output_cm_scores_(pred, np.transpose(test_labels))
```

```
In [1]: def train_SVM(hog_img, labels, fold_index):
        SVM=svm.LinearSVC(random_state=0)
        mask=get_split_mask(hog_img)
        train_hog, train_labels, test_hog, test_labels=get_split_data(images=hog_img, labels=labels, mask=mask)
        SVM.fit(train_hog, train_labels)
        pred=SVM.predict(test_hog)
        output_cm_scores_(pred,np.transpose(test_labels))
```

## 2 Training Results:

### 2.0.1 Sigmoid Activation: Fold - 1 , Learning Rate=1e-3

```
In [71]: model0=MLP(input_size, h1_size, h2_size, output_size)
        train(images=hog_img, labels=labels, epochs=10, fold_index=0, model=model0)
```


```
Epoch 1 Iteration 1 : Avg Loss = 2.5237310604414125
Epoch 1 Iteration 201 : Avg Loss = 0.884445213530523
Epoch 1 Iteration 401 : Avg Loss = 0.35903156786276985
Epoch 1 Iteration 601 : Avg Loss = 0.3655924703778429
Epoch 2 Iteration 1 : Avg Loss = 0.10902911774288424
Epoch 2 Iteration 201 : Avg Loss = 0.36398440514331526
Epoch 2 Iteration 401 : Avg Loss = 0.2258086851886312
Epoch 2 Iteration 601 : Avg Loss = 0.22260797401805532
Epoch 3 Iteration 1 : Avg Loss = 0.05129968587343006
Epoch 3 Iteration 201 : Avg Loss = 0.26492773978105
Epoch 3 Iteration 401 : Avg Loss = 0.1559271781710433
Epoch 3 Iteration 601 : Avg Loss = 0.1737974770920139
Epoch 4 Iteration 1 : Avg Loss = 0.03843904655200772
Epoch 4 Iteration 201 : Avg Loss = 0.22684122539521784
Epoch 4 Iteration 401 : Avg Loss = 0.11652301922279068
Epoch 4 Iteration 601 : Avg Loss = 0.14757747366434396
Epoch 5 Iteration 1 : Avg Loss = 0.031926515058978984
Epoch 5 Iteration 201 : Avg Loss = 0.20134238224677115
Epoch 5 Iteration 401 : Avg Loss = 0.0924885450678404
Epoch 5 Iteration 601 : Avg Loss = 0.12984148966792264
Epoch 6 Iteration 1 : Avg Loss = 0.028390987012759175
Epoch 6 Iteration 201 : Avg Loss = 0.18114311356343926
Epoch 6 Iteration 401 : Avg Loss = 0.0749974610359073
Epoch 6 Iteration 601 : Avg Loss = 0.11802888229472246
Epoch 7 Iteration 1 : Avg Loss = 0.02568164330554715
Epoch 7 Iteration 201 : Avg Loss = 0.16530945625132126
Epoch 7 Iteration 401 : Avg Loss = 0.062439104136879364
Epoch 7 Iteration 601 : Avg Loss = 0.10933906205265685
Epoch 8 Iteration 1 : Avg Loss = 0.02327578612019556
Epoch 8 Iteration 201 : Avg Loss = 0.15203635959131023
Epoch 8 Iteration 401 : Avg Loss = 0.053729359559740106
Epoch 8 Iteration 601 : Avg Loss = 0.10221266638664009
Epoch 9 Iteration 1 : Avg Loss = 0.02109309421811379
```



Epoch 9 Iteration 201 : Avg Loss = 0.14007147946780688  
 Epoch 9 Iteration 401 : Avg Loss = 0.04746798105823839  
 Epoch 9 Iteration 601 : Avg Loss = 0.0961481985689468  
 Epoch 10 Iteration 1 : Avg Loss = 0.018899770552620404  
 Epoch 10 Iteration 201 : Avg Loss = 0.12860275403795643  
 Epoch 10 Iteration 401 : Avg Loss = 0.04269410014086432  
 Epoch 10 Iteration 601 : Avg Loss = 0.09044278243873298  
 40



Accuracy = 0.96625  
 Precision = [0.98682043 0.98937785 0.95641447 0.94204019 0.98188406 0.97920605  
 0.96700508 0.96642097 0.9691358 0.92851469]  
 Recall = [0.98763397 0.98415094 0.97485331 0.97833066 0.94838145 0.95308188  
 0.99046794 0.94551282 0.92508418 0.97173732]  
 F1 Score = [0.98722703 0.98675747 0.96554587 0.95984252 0.96484201 0.96596737  
 0.97859589 0.95585257 0.94659776 0.94963444]



## 2.0.2 Sigmoid Activation: Fold - 2 , Learning Rate=1e-3


```
In [73]: model1=MLP(input_size, h1_size, h2_size, output_size)
         train(images=hog_img, labels=labels, epochs=10, fold_index=1, model=model1)
```

```
Epoch 1 Iteration 1 : Avg Loss = 2.3486130349224137
Epoch 1 Iteration 201 : Avg Loss = 0.8805077968073363
Epoch 1 Iteration 401 : Avg Loss = 0.35118200690820534
Epoch 1 Iteration 601 : Avg Loss = 0.3774719906840963
Epoch 2 Iteration 1 : Avg Loss = 0.15647418099205188
Epoch 2 Iteration 201 : Avg Loss = 0.33701204427867015
Epoch 2 Iteration 401 : Avg Loss = 0.2292529184510146
Epoch 2 Iteration 601 : Avg Loss = 0.2673917212471495
Epoch 3 Iteration 1 : Avg Loss = 0.10403024286388032
Epoch 3 Iteration 201 : Avg Loss = 0.2429156371271068
Epoch 3 Iteration 401 : Avg Loss = 0.17408084793344306
Epoch 3 Iteration 601 : Avg Loss = 0.2241918291408031
```

Epoch 4 Iteration 1 : Avg Loss = 0.08549408312668665  
Epoch 4 Iteration 201 : Avg Loss = 0.1959961144193908  
Epoch 4 Iteration 401 : Avg Loss = 0.14028176742571927  
Epoch 4 Iteration 601 : Avg Loss = 0.1952009737202599  
Epoch 5 Iteration 1 : Avg Loss = 0.07061118638707099  
Epoch 5 Iteration 201 : Avg Loss = 0.1664451775125741  
Epoch 5 Iteration 401 : Avg Loss = 0.12167676985414602  
Epoch 5 Iteration 601 : Avg Loss = 0.17237019309312093  
Epoch 6 Iteration 1 : Avg Loss = 0.05740647102003866  
Epoch 6 Iteration 201 : Avg Loss = 0.14803157482959878  
Epoch 6 Iteration 401 : Avg Loss = 0.10812726609504619  
Epoch 6 Iteration 601 : Avg Loss = 0.1543118751711875  
Epoch 7 Iteration 1 : Avg Loss = 0.04705788424598002  
Epoch 7 Iteration 201 : Avg Loss = 0.13573084250629824  
Epoch 7 Iteration 401 : Avg Loss = 0.09490796121066841  
Epoch 7 Iteration 601 : Avg Loss = 0.1405772328128907  
Epoch 8 Iteration 1 : Avg Loss = 0.040060624910670015  
Epoch 8 Iteration 201 : Avg Loss = 0.12595932838728982  
Epoch 8 Iteration 401 : Avg Loss = 0.08337007569707795  
Epoch 8 Iteration 601 : Avg Loss = 0.1310908655700931  
Epoch 9 Iteration 1 : Avg Loss = 0.03522286143880078  
Epoch 9 Iteration 201 : Avg Loss = 0.11782556391561515  
Epoch 9 Iteration 401 : Avg Loss = 0.07476002773884541  
Epoch 9 Iteration 601 : Avg Loss = 0.12591232518534026  
Epoch 10 Iteration 1 : Avg Loss = 0.031476027991174875  
Epoch 10 Iteration 201 : Avg Loss = 0.11150040934074947  
Epoch 10 Iteration 401 : Avg Loss = 0.06873461115827453  
Epoch 10 Iteration 601 : Avg Loss = 0.12374684228922815



```
Accuracy = 0.9668333333333333
Precision = [0.98489933 0.98317484 0.96714406 0.93044822 0.98196573 0.9820949
0.96473552 0.97442244 0.97563177 0.92804579]
Recall = [0.98160535 0.98606016 0.96632997 0.97806661 0.93960311 0.9580786
0.98966408 0.95938262 0.93512111 0.97008547]
F1 Score = [0.98324958 0.98461538 0.96673684 0.95366337 0.96031746 0.96993811
0.97704082 0.96684404 0.954947 0.94860008]
```




### 2.0.3 Sigmoid Activation: Fold - 3 , Learning Rate=1e-3

```
In [74]: model2=MLP(input_size, h1_size, h2_size, output_size)
         train(images=hog_img, labels=labels, epochs=10, fold_index=2, model=model2)
```

```
Epoch 1 Iteration 1 : Avg Loss = 2.5608770701074492
Epoch 1 Iteration 201 : Avg Loss = 0.7373994279409963
Epoch 1 Iteration 401 : Avg Loss = 0.3313299908022806
Epoch 1 Iteration 601 : Avg Loss = 0.36246301735684483
Epoch 2 Iteration 1 : Avg Loss = 0.15393990487147974
Epoch 2 Iteration 201 : Avg Loss = 0.19700579478168634
Epoch 2 Iteration 401 : Avg Loss = 0.23658798412097798
Epoch 2 Iteration 601 : Avg Loss = 0.2266668083979541
Epoch 3 Iteration 1 : Avg Loss = 0.10315344828157831
Epoch 3 Iteration 201 : Avg Loss = 0.16379690605442782
Epoch 3 Iteration 401 : Avg Loss = 0.1771364006606284
Epoch 3 Iteration 601 : Avg Loss = 0.16940896725257232
```

Epoch 4 Iteration 1 : Avg Loss = 0.08072841023500066  
Epoch 4 Iteration 201 : Avg Loss = 0.14181660492921588  
Epoch 4 Iteration 401 : Avg Loss = 0.14119408430178249  
Epoch 4 Iteration 601 : Avg Loss = 0.1442354279861858  
Epoch 5 Iteration 1 : Avg Loss = 0.0663022074929781  
Epoch 5 Iteration 201 : Avg Loss = 0.12267082325345313  
Epoch 5 Iteration 401 : Avg Loss = 0.11731573477474488  
Epoch 5 Iteration 601 : Avg Loss = 0.12879216723723655  
Epoch 6 Iteration 1 : Avg Loss = 0.055904690064533566  
Epoch 6 Iteration 201 : Avg Loss = 0.1071905701190896  
Epoch 6 Iteration 401 : Avg Loss = 0.09837409225715735  
Epoch 6 Iteration 601 : Avg Loss = 0.11958445332605273  
Epoch 7 Iteration 1 : Avg Loss = 0.047152172222009896  
Epoch 7 Iteration 201 : Avg Loss = 0.09413525012848532  
Epoch 7 Iteration 401 : Avg Loss = 0.08372216596109702  
Epoch 7 Iteration 601 : Avg Loss = 0.11412888297040102  
Epoch 8 Iteration 1 : Avg Loss = 0.04019930864665801  
Epoch 8 Iteration 201 : Avg Loss = 0.08303096591017846  
Epoch 8 Iteration 401 : Avg Loss = 0.07303894342702852  
Epoch 8 Iteration 601 : Avg Loss = 0.11105829925024217  
Epoch 9 Iteration 1 : Avg Loss = 0.03505150696744377  
Epoch 9 Iteration 201 : Avg Loss = 0.07428206722388714  
Epoch 9 Iteration 401 : Avg Loss = 0.06519559641680223  
Epoch 9 Iteration 601 : Avg Loss = 0.10938989614345987  
Epoch 10 Iteration 1 : Avg Loss = 0.031332867444667886  
Epoch 10 Iteration 201 : Avg Loss = 0.0681132447318449  
Epoch 10 Iteration 401 : Avg Loss = 0.058896756722482776  
Epoch 10 Iteration 601 : Avg Loss = 0.1076279856585364




Accuracy = 0.9691666666666666

Precision = [0.99087894 0.98875562 0.96545914 0.94784026 0.98577778 0.97619048  
0.96523848 0.98146656 0.97610619 0.91686461]

Recall = [0.9859736 0.98359433 0.97614991 0.97240803 0.94624573 0.9743346  
0.995 0.95379796 0.92301255 0.98052498]

F1 Score = [0.98842018 0.98616822 0.9707751 0.95996698 0.96560731 0.97526166  
0.97989331 0.96743447 0.9488172 0.94762684]




## 2.0.4 Sigmoid Activation: Fold - 4 , Learning Rate=1e-3

```
In [75]: model3=MLP(input_size, h1_size, h2_size, output_size)
         train(images=hog_img, labels=labels, epochs=10, fold_index=3, model=model3)
```

```
Epoch 1 Iteration 1 : Avg Loss = 2.61727318186808
Epoch 1 Iteration 201 : Avg Loss = 0.8704242156675674
Epoch 1 Iteration 401 : Avg Loss = 0.3838612418083182
Epoch 1 Iteration 601 : Avg Loss = 0.34807333461493795
Epoch 2 Iteration 1 : Avg Loss = 0.1340971846288544
Epoch 2 Iteration 201 : Avg Loss = 0.21459745950758458
Epoch 2 Iteration 401 : Avg Loss = 0.2222683176918635
Epoch 2 Iteration 601 : Avg Loss = 0.2132056718473819
Epoch 3 Iteration 1 : Avg Loss = 0.08956229127077667
Epoch 3 Iteration 201 : Avg Loss = 0.17164518048850474
Epoch 3 Iteration 401 : Avg Loss = 0.16195204657385218
Epoch 3 Iteration 601 : Avg Loss = 0.1658042036857238
```



Epoch 4 Iteration 1 : Avg Loss = 0.07350707623568369  
Epoch 4 Iteration 201 : Avg Loss = 0.1488588781812995  
Epoch 4 Iteration 401 : Avg Loss = 0.13192430967350516  
Epoch 4 Iteration 601 : Avg Loss = 0.14221631945540786  
Epoch 5 Iteration 1 : Avg Loss = 0.06533641702178208  
Epoch 5 Iteration 201 : Avg Loss = 0.12490333848970356  
Epoch 5 Iteration 401 : Avg Loss = 0.11246842146802903  
Epoch 5 Iteration 601 : Avg Loss = 0.12876281625083752  
Epoch 6 Iteration 1 : Avg Loss = 0.06024927838476302  
Epoch 6 Iteration 201 : Avg Loss = 0.1052212343898589  
Epoch 6 Iteration 401 : Avg Loss = 0.09660097093705501  
Epoch 6 Iteration 601 : Avg Loss = 0.12060765625201145  
Epoch 7 Iteration 1 : Avg Loss = 0.056257417432786436  
Epoch 7 Iteration 201 : Avg Loss = 0.09054073301065518  
Epoch 7 Iteration 401 : Avg Loss = 0.08167994378466077  
Epoch 7 Iteration 601 : Avg Loss = 0.1141990344057979  
Epoch 8 Iteration 1 : Avg Loss = 0.05274694745301829  
Epoch 8 Iteration 201 : Avg Loss = 0.07935088781530081  
Epoch 8 Iteration 401 : Avg Loss = 0.0677111509683827  
Epoch 8 Iteration 601 : Avg Loss = 0.10807128947816369  
Epoch 9 Iteration 1 : Avg Loss = 0.049575219308992016  
Epoch 9 Iteration 201 : Avg Loss = 0.07078636261030061  
Epoch 9 Iteration 401 : Avg Loss = 0.05549920232938613  
Epoch 9 Iteration 601 : Avg Loss = 0.10155088564853891  
Epoch 10 Iteration 1 : Avg Loss = 0.04644923545298721  
Epoch 10 Iteration 201 : Avg Loss = 0.06440703772560766  
Epoch 10 Iteration 401 : Avg Loss = 0.046037476503504635  
Epoch 10 Iteration 601 : Avg Loss = 0.09473484911361188




Accuracy = 0.9705

Precision = [0.97584124 0.98529412 0.95929769 0.9625498 0.99021352 0.98167792  
0.97233523 0.96024942 0.97971781 0.94246353]

Recall = [0.98950131 0.98301158 0.9796251 0.97734628 0.93529412 0.96128423  
0.99087894 0.96779262 0.94312394 0.97485331]

F1 Score = [0.98262381 0.98415153 0.96935484 0.96989161 0.96197061 0.97137405  
0.98151951 0.96400626 0.96107266 0.95838484]




## 2.0.5 ReLU Activation: Fold - 1 , Learning Rate=5e-4

```
In [78]: model0=MLP(input_size, h1_size, h2_size, output_size)
         train(images=hog_img, labels=labels, epochs=10, fold_index=0, model=model0)
```

```
Epoch 1 Iteration 1 : Avg Loss = 9.120454114316132
Epoch 1 Iteration 201 : Avg Loss = 0.3695540096432324
Epoch 1 Iteration 401 : Avg Loss = 0.2563435714996035
Epoch 1 Iteration 601 : Avg Loss = 0.2329868392901371
Epoch 2 Iteration 1 : Avg Loss = 0.033891931897472014
Epoch 2 Iteration 201 : Avg Loss = 0.32103634967947414
Epoch 2 Iteration 401 : Avg Loss = 0.16219619921774103
Epoch 2 Iteration 601 : Avg Loss = 0.13502726609932056
Epoch 3 Iteration 1 : Avg Loss = 0.01103076294020759
Epoch 3 Iteration 201 : Avg Loss = 0.2586787284343916
Epoch 3 Iteration 401 : Avg Loss = 0.1622301980850681
Epoch 3 Iteration 601 : Avg Loss = 0.0983698625045581
```

Epoch 4 Iteration 1 : Avg Loss = 0.008943218464390768  
Epoch 4 Iteration 201 : Avg Loss = 0.2421379000804832  
Epoch 4 Iteration 401 : Avg Loss = 0.16164700315441405  
Epoch 4 Iteration 601 : Avg Loss = 0.15484803320905666  
Epoch 5 Iteration 1 : Avg Loss = 0.009118390863813414  
Epoch 5 Iteration 201 : Avg Loss = 0.21582579941580404  
Epoch 5 Iteration 401 : Avg Loss = 0.15499095557923986  
Epoch 5 Iteration 601 : Avg Loss = 0.09493966806324963  
Epoch 6 Iteration 1 : Avg Loss = 0.010490351099428197  
Epoch 6 Iteration 201 : Avg Loss = 0.18654996837916796  
Epoch 6 Iteration 401 : Avg Loss = 0.09390297338679306  
Epoch 6 Iteration 601 : Avg Loss = 0.046540519319866865  
Epoch 7 Iteration 1 : Avg Loss = 0.03453948600237958  
Epoch 7 Iteration 201 : Avg Loss = 0.16441314588527928  
Epoch 7 Iteration 401 : Avg Loss = 0.0819965470904304  
Epoch 7 Iteration 601 : Avg Loss = 0.032412598727772135  
Epoch 8 Iteration 1 : Avg Loss = 0.021400701645603367  
Epoch 8 Iteration 201 : Avg Loss = 0.18457018942571835  
Epoch 8 Iteration 401 : Avg Loss = 0.05032321096995107  
Epoch 8 Iteration 601 : Avg Loss = 0.015144899475195855  
Epoch 9 Iteration 1 : Avg Loss = 0.006296161301422975  
Epoch 9 Iteration 201 : Avg Loss = 0.17762634986350748  
Epoch 9 Iteration 401 : Avg Loss = 0.08052973097848583  
Epoch 9 Iteration 601 : Avg Loss = 0.03981868745681544  
Epoch 10 Iteration 1 : Avg Loss = 0.007603637702699635  
Epoch 10 Iteration 201 : Avg Loss = 0.10729036699606607  
Epoch 10 Iteration 401 : Avg Loss = 0.03250305093712224  
Epoch 10 Iteration 601 : Avg Loss = 0.0492099029442807




Accuracy = 0.96825

Precision = [0.98676592 0.98429319 0.95358306 0.95070423 0.96542783 0.98745174  
0.96537162 0.96593674 0.95768566 0.96694915]

Recall = [0.98351195 0.99320755 0.98155909 0.97512039 0.97725284 0.94112236  
0.99046794 0.95432692 0.93350168 0.94846218]

F1 Score = [0.98513625 0.98873028 0.96736886 0.96275753 0.97130435 0.96373057  
0.97775877 0.96009674 0.94543905 0.95761645]




## 2.0.6 ReLU Activation: Fold - 2 , Learning Rate=5e-4

```
In [79]: model1=MLP(input_size, h1_size, h2_size, output_size)
         train(images=hog_img, labels=labels, epochs=10, fold_index=1, model=model1)
```

```
Epoch 1 Iteration 1 : Avg Loss = 5.686153496418004
Epoch 1 Iteration 201 : Avg Loss = 0.4629774230970316
Epoch 1 Iteration 401 : Avg Loss = 0.17903476425572615
Epoch 1 Iteration 601 : Avg Loss = 0.21012480867141517
Epoch 2 Iteration 1 : Avg Loss = 0.05348664507058814
Epoch 2 Iteration 201 : Avg Loss = 0.22230705885393404
Epoch 2 Iteration 401 : Avg Loss = 0.1792182683996405
Epoch 2 Iteration 601 : Avg Loss = 0.11641789411432661
Epoch 3 Iteration 1 : Avg Loss = 0.07433463501481997
Epoch 3 Iteration 201 : Avg Loss = 0.16261710529705697
Epoch 3 Iteration 401 : Avg Loss = 0.12071240316830212
Epoch 3 Iteration 601 : Avg Loss = 0.13939150151652466
```

Epoch 4 Iteration 1 : Avg Loss = 0.052717498876482194  
Epoch 4 Iteration 201 : Avg Loss = 0.13673239208360083  
Epoch 4 Iteration 401 : Avg Loss = 0.07667629642699361  
Epoch 4 Iteration 601 : Avg Loss = 0.08814960443377175  
Epoch 5 Iteration 1 : Avg Loss = 0.02887043707787417  
Epoch 5 Iteration 201 : Avg Loss = 0.09594533241439655  
Epoch 5 Iteration 401 : Avg Loss = 0.09157138607078297  
Epoch 5 Iteration 601 : Avg Loss = 0.11366887943828569  
Epoch 6 Iteration 1 : Avg Loss = 0.034516186806697255  
Epoch 6 Iteration 201 : Avg Loss = 0.10210740009288341  
Epoch 6 Iteration 401 : Avg Loss = 0.06621207463419  
Epoch 6 Iteration 601 : Avg Loss = 0.08332966772978152  
Epoch 7 Iteration 1 : Avg Loss = 0.021335071446704185  
Epoch 7 Iteration 201 : Avg Loss = 0.09011580377665288  
Epoch 7 Iteration 401 : Avg Loss = 0.05149619979238955  
Epoch 7 Iteration 601 : Avg Loss = 0.05902780063687792  
Epoch 8 Iteration 1 : Avg Loss = 0.015797329020319446  
Epoch 8 Iteration 201 : Avg Loss = 0.13103298408229472  
Epoch 8 Iteration 401 : Avg Loss = 0.0426423078357665  
Epoch 8 Iteration 601 : Avg Loss = 0.07821632826994868  
Epoch 9 Iteration 1 : Avg Loss = 0.028729327741394466  
Epoch 9 Iteration 201 : Avg Loss = 0.07311715606131325  
Epoch 9 Iteration 401 : Avg Loss = 0.057211775680836345  
Epoch 9 Iteration 601 : Avg Loss = 0.019621590543342525  
Epoch 10 Iteration 1 : Avg Loss = 0.015075036365580278  
Epoch 10 Iteration 201 : Avg Loss = 0.04951006118409535  
Epoch 10 Iteration 401 : Avg Loss = 0.028911340485648574  
Epoch 10 Iteration 601 : Avg Loss = 0.015299236908133037




Accuracy = 0.9673333333333334

Precision = [0.98826488 0.98243045 0.97140454 0.95987159 0.97942755 0.98454545  
0.97198642 0.9814346 0.95682211 0.90134175]

Recall = [0.98578595 0.98459281 0.97222222 0.97156783 0.94477998 0.94585153  
0.98621878 0.94476036 0.95847751 0.97606838]

F1 Score = [0.98702386 0.98351044 0.97181321 0.9656843 0.96179183 0.96481069  
0.97905088 0.96274834 0.95764909 0.93721789]






## 2.0.7 ReLU Activation: Fold - 3 , Learning Rate=5e-4

```
In [80]: model2=MLP(input_size, h1_size, h2_size, output_size)
         train(images=hog_img, labels=labels, epochs=10, fold_index=2, model=model2)
```

```
Epoch 1 Iteration 1 : Avg Loss = 7.7289568417855055
Epoch 1 Iteration 201 : Avg Loss = 0.4767908575414811
Epoch 1 Iteration 401 : Avg Loss = 0.3539230563369344
Epoch 1 Iteration 601 : Avg Loss = 0.19498260146854265
Epoch 2 Iteration 1 : Avg Loss = 0.11988889014426031
Epoch 2 Iteration 201 : Avg Loss = 0.15201172364305393
Epoch 2 Iteration 401 : Avg Loss = 0.1918882570027483
Epoch 2 Iteration 601 : Avg Loss = 0.1381031830695138
Epoch 3 Iteration 1 : Avg Loss = 0.10103946291874161
Epoch 3 Iteration 201 : Avg Loss = 0.10967983019753948
Epoch 3 Iteration 401 : Avg Loss = 0.09412550603307032
Epoch 3 Iteration 601 : Avg Loss = 0.08914121278336554
```


Epoch 4 Iteration 1 : Avg Loss = 0.07045661571137758  
Epoch 4 Iteration 201 : Avg Loss = 0.07562473876215411  
Epoch 4 Iteration 401 : Avg Loss = 0.0635467723364539  
Epoch 4 Iteration 601 : Avg Loss = 0.07559270580111976  
Epoch 5 Iteration 1 : Avg Loss = 0.04826412348498155  
Epoch 5 Iteration 201 : Avg Loss = 0.05572940869537378  
Epoch 5 Iteration 401 : Avg Loss = 0.04486971997321076  
Epoch 5 Iteration 601 : Avg Loss = 0.07322558322903758  
Epoch 6 Iteration 1 : Avg Loss = 0.0504528623033276  
Epoch 6 Iteration 201 : Avg Loss = 0.0810150363080389  
Epoch 6 Iteration 401 : Avg Loss = 0.051240136667419355  
Epoch 6 Iteration 601 : Avg Loss = 0.06541453024892892  
Epoch 7 Iteration 1 : Avg Loss = 0.028444536923683754  
Epoch 7 Iteration 201 : Avg Loss = 0.03727082052452081  
Epoch 7 Iteration 401 : Avg Loss = 0.030093062872504614  
Epoch 7 Iteration 601 : Avg Loss = 0.06203271373423305  
Epoch 8 Iteration 1 : Avg Loss = 0.024002034205703626  
Epoch 8 Iteration 201 : Avg Loss = 0.03248176947002494  
Epoch 8 Iteration 401 : Avg Loss = 0.03921060631717962  
Epoch 8 Iteration 601 : Avg Loss = 0.052338320996058875  
Epoch 9 Iteration 1 : Avg Loss = 0.03871390241072076  
Epoch 9 Iteration 201 : Avg Loss = 0.05260957846975152  
Epoch 9 Iteration 401 : Avg Loss = 0.03800104195705883  
Epoch 9 Iteration 601 : Avg Loss = 0.02829026006163452  
Epoch 10 Iteration 1 : Avg Loss = 0.03762736967892924  
Epoch 10 Iteration 201 : Avg Loss = 0.022726322083105755  
Epoch 10 Iteration 401 : Avg Loss = 0.025565760971025427  
Epoch 10 Iteration 601 : Avg Loss = 0.02957233267804254



```

Accuracy = 0.9659166666666666
Precision = [0.98763397 0.98805078 0.96711636 0.94809408 0.98405669 0.98336595
0.95458167 0.97502082 0.97465035 0.9029734 ]
Recall = [0.98844884 0.98657718 0.9770017 0.97742475 0.94795222 0.95532319
0.99833333 0.91699295 0.93305439 0.97713802]
F1 Score = [0.98804124 0.98731343 0.9720339 0.96253602 0.9656671 0.96914176
0.97596741 0.94511703 0.95339889 0.93859292]

```




## 2.0.8 ReLU Activation: Fold - 4 , Learning Rate=5e-4


```
In [81]: model3=MLP(input_size, h1_size, h2_size, output_size)
         train(images=hog_img, labels=labels, epochs=10, fold_index=3, model=model3)
```

```
Epoch 1 Iteration 1 : Avg Loss = 8.73148664539806
Epoch 1 Iteration 201 : Avg Loss = 0.3428013177751035
Epoch 1 Iteration 401 : Avg Loss = 0.23369188384478579
Epoch 1 Iteration 601 : Avg Loss = 0.2042677366311289
Epoch 2 Iteration 1 : Avg Loss = 0.17110737835237658
Epoch 2 Iteration 201 : Avg Loss = 0.20369544499754108
Epoch 2 Iteration 401 : Avg Loss = 0.17257317536873873
Epoch 2 Iteration 601 : Avg Loss = 0.13164172663356488
Epoch 3 Iteration 1 : Avg Loss = 0.08116140438953583
Epoch 3 Iteration 201 : Avg Loss = 0.15247239997726222
Epoch 3 Iteration 401 : Avg Loss = 0.10891679864831445
Epoch 3 Iteration 601 : Avg Loss = 0.0962537658825236
```

Epoch 4 Iteration 1 : Avg Loss = 0.0628560789125494  
Epoch 4 Iteration 201 : Avg Loss = 0.13475861205303025  
Epoch 4 Iteration 401 : Avg Loss = 0.08095876305624854  
Epoch 4 Iteration 601 : Avg Loss = 0.06913590105589297  
Epoch 5 Iteration 1 : Avg Loss = 0.06665181353210972  
Epoch 5 Iteration 201 : Avg Loss = 0.0662147601957997  
Epoch 5 Iteration 401 : Avg Loss = 0.04630521229031184  
Epoch 5 Iteration 601 : Avg Loss = 0.05317881752069942  
Epoch 6 Iteration 1 : Avg Loss = 0.03444957516862797  
Epoch 6 Iteration 201 : Avg Loss = 0.053649075961792816  
Epoch 6 Iteration 401 : Avg Loss = 0.03012052388546444  
Epoch 6 Iteration 601 : Avg Loss = 0.043634725720439915  
Epoch 7 Iteration 1 : Avg Loss = 0.031534681066791005  
Epoch 7 Iteration 201 : Avg Loss = 0.058949691252450684  
Epoch 7 Iteration 401 : Avg Loss = 0.023746311551862208  
Epoch 7 Iteration 601 : Avg Loss = 0.02947613191008299  
Epoch 8 Iteration 1 : Avg Loss = 0.028301773012016883  
Epoch 8 Iteration 201 : Avg Loss = 0.03299403901221629  
Epoch 8 Iteration 401 : Avg Loss = 0.008379397744295641  
Epoch 8 Iteration 601 : Avg Loss = 0.06611256263442201  
Epoch 9 Iteration 1 : Avg Loss = 0.0140484036381062  
Epoch 9 Iteration 201 : Avg Loss = 0.05860864870432068  
Epoch 9 Iteration 401 : Avg Loss = 0.016484641139775404  
Epoch 9 Iteration 601 : Avg Loss = 0.05300744201430871  
Epoch 10 Iteration 1 : Avg Loss = 0.017310534388658882  
Epoch 10 Iteration 201 : Avg Loss = 0.02800197666738998  
Epoch 10 Iteration 401 : Avg Loss = 0.0027657788361300883  
Epoch 10 Iteration 601 : Avg Loss = 0.032562322922983486



Accuracy = 0.9680833333333333  
Precision = [0.98688811 0.98684211 0.95863166 0.95800317 0.99539595 0.9892051  
0.97547016 0.96075353 0.94920899 0.92834138]  
Recall = [0.98775153 0.98455598 0.98207009 0.97815534 0.90840336 0.95184136  
0.98922056 0.96150825 0.96774194 0.96647108]  
F1 Score = [0.98731963 0.98569772 0.97020934 0.96797438 0.94991213 0.97016362  
0.98229724 0.96113074 0.95838588 0.94702259]




## 2.0.9 ReLU Activation: Fold - 5 , Learning Rate=5e-4

```
In [82]: model4=MLP(input_size, h1_size, h2_size, output_size)
         train(images=hog_img, labels=labels, epochs=10, fold_index=4, model=model4)
```

```
Epoch 1 Iteration 1 : Avg Loss = 6.452419998397894
Epoch 1 Iteration 201 : Avg Loss = 0.40023061029722135
Epoch 1 Iteration 401 : Avg Loss = 0.17510700787350808
Epoch 1 Iteration 601 : Avg Loss = 0.11678254612431363
Epoch 2 Iteration 1 : Avg Loss = 0.07297036524886916
Epoch 2 Iteration 201 : Avg Loss = 0.1881079525044066
Epoch 2 Iteration 401 : Avg Loss = 0.1888615387786949
Epoch 2 Iteration 601 : Avg Loss = 0.14828309971065878
Epoch 3 Iteration 1 : Avg Loss = 0.0691296822765721
Epoch 3 Iteration 201 : Avg Loss = 0.10173555418989014
Epoch 3 Iteration 401 : Avg Loss = 0.14566269985428876
Epoch 3 Iteration 601 : Avg Loss = 0.11688379121844969
```

Epoch 4 Iteration 1 : Avg Loss = 0.04856110558598667  
Epoch 4 Iteration 201 : Avg Loss = 0.07830608669722441  
Epoch 4 Iteration 401 : Avg Loss = 0.07990529918493103  
Epoch 4 Iteration 601 : Avg Loss = 0.09795277032129626  
Epoch 5 Iteration 1 : Avg Loss = 0.028402865971453164  
Epoch 5 Iteration 201 : Avg Loss = 0.08302671915807842  
Epoch 5 Iteration 401 : Avg Loss = 0.038439255179968365  
Epoch 5 Iteration 601 : Avg Loss = 0.04206275166613032  
Epoch 6 Iteration 1 : Avg Loss = 0.024634110359063113  
Epoch 6 Iteration 201 : Avg Loss = 0.0641232264861061  
Epoch 6 Iteration 401 : Avg Loss = 0.03140132256054211  
Epoch 6 Iteration 601 : Avg Loss = 0.04534162271382405  
Epoch 7 Iteration 1 : Avg Loss = 0.02920686127899648  
Epoch 7 Iteration 201 : Avg Loss = 0.04026503338404281  
Epoch 7 Iteration 401 : Avg Loss = 0.02642523893575877  
Epoch 7 Iteration 601 : Avg Loss = 0.056157282918288504  
Epoch 8 Iteration 1 : Avg Loss = 0.05163475866612013  
Epoch 8 Iteration 201 : Avg Loss = 0.05030746820019917  
Epoch 8 Iteration 401 : Avg Loss = 0.012630129906847025  
Epoch 8 Iteration 601 : Avg Loss = 0.02424433438057641  
Epoch 9 Iteration 1 : Avg Loss = 0.03185902907025416  
Epoch 9 Iteration 201 : Avg Loss = 0.03967608046544924  
Epoch 9 Iteration 401 : Avg Loss = 0.023369219452937105  
Epoch 9 Iteration 601 : Avg Loss = 0.014198182768065767  
Epoch 10 Iteration 1 : Avg Loss = 0.014902764776699171  
Epoch 10 Iteration 201 : Avg Loss = 0.062476652899227415  
Epoch 10 Iteration 401 : Avg Loss = 0.014880773452666824  
Epoch 10 Iteration 601 : Avg Loss = 0.016509734007795866






Accuracy = 0.97075

Precision = [0.98040886 0.98308668 0.9734589 0.94976077 0.97108844 0.95687332  
0.99146758 0.96209677 0.96024735 0.97743056]

Recall = [0.9930975 0.98377997 0.96683673 0.97463175 0.96943973 0.98794063  
0.97076023 0.96521036 0.95855379 0.93677205]

F1 Score = [0.98671239 0.9834332 0.97013652 0.96203554 0.97026338 0.97215883  
0.98100464 0.96365105 0.95939982 0.9566695 ]



## 2.0.10 KNN: Fold-1


In [117]: train\_KNN(hog\_img, labels, fold\_index=0)

Accuracy = 0.9524166666666667

Precision = [0.95335968 0.99086062 0.97929249 0.93665875 0.95259799 0.9804878  
0.96872358 0.95496247 0.91451613 0.89888535]

Recall = [0.99422918 0.98188679 0.95138307 0.9494382 0.91426072 0.92456302  
0.99306759 0.91746795 0.95454545 0.93848712]

F1 Score = [0.97336562 0.9863533 0.96513605 0.94300518 0.93303571 0.95170455  
0.98074454 0.9358398 0.93410214 0.91825946]



## 2.0.11 KNN: Fold-2


In [120]: `train_KNN(hog_img, labels, fold_index=1)`

Accuracy = 0.96475

Precision = [0.97532895 0.98965262 0.98289136 0.94732642 0.98001817 0.98903108  
0.97292724 0.97145256 0.93366501 0.90850202]

Recall = [0.9916388 0.98239178 0.96717172 0.9642567 0.93097498 0.94497817  
0.99052541 0.93988627 0.97404844 0.95897436]

F1 Score = [0.98341625 0.98600884 0.97496818 0.95571659 0.95486726 0.9665029  
0.98164746 0.95540875 0.9534293 0.93305613]



## 2.0.12 KNN: Fold-3


In [19]: `train_KNN(hog_img, labels, fold_index=2)`

Accuracy = 0.9575833333333333

Precision = [0.95833333 0.98688047 0.97846684 0.94741167 0.96930946 0.98510427  
0.9672696 0.9734589 0.93021412 0.88880368]

Recall = [0.99207746 0.97973951 0.95542473 0.94975288 0.92892157 0.9271028  
0.99468556 0.92288961 0.96622735 0.95785124]

F1 Score = [0.97491349 0.98329702 0.96680851 0.94858083 0.94868586 0.95522388  
0.98078603 0.9475 0.94787879 0.9220366 ]



### 2.0.13 KNN: Fold-4


In [21]: `train_KNN(hog_img, labels, fold_index=3)`

Accuracy = 0.95475

Precision = [0.9699187 0.98171178 0.98198971 0.92598684 0.96744186 0.9782197  
0.96405751 0.97729184 0.92460646 0.88003221]

Recall = [0.99333888 0.9874908 0.94706369 0.94781145 0.91148116 0.93146979  
0.98934426 0.92222222 0.9530316 0.95709282]

F1 Score = [0.9814891 0.98459281 0.96421053 0.93677205 0.93862816 0.95427252  
0.97653722 0.94895876 0.93860387 0.91694631]



## 2.0.14 KNN: Fold-5


In [23]: `train_KNN(hog_img, labels, fold_index=4)`

Accuracy = 0.9555833333333333

Precision = [0.9645799 0.98604651 0.97719595 0.92211838 0.96332737 0.97965116  
0.9693301 0.96522477 0.9343696 0.90168971]

Recall = [0.99237288 0.98223938 0.95540875 0.94947875 0.93246753 0.95557656  
0.98766447 0.91405622 0.94250871 0.94297189]

F1 Score = [0.97827903 0.98413926 0.96617954 0.93559858 0.94764628 0.96746411  
0.97841141 0.93894389 0.93842151 0.92186887]



## 2.0.15 SVM: Fold-1


In [14]: `train_SVM(hog_img, labels, 0)`

Accuracy = 0.9051666666666667

Precision = [0.95993459 0.96594203 0.79044118 0.98103267 0.86525974 0.89948007  
0.96822742 0.90025795 0.83442623 0.91006233]

Recall = [0.97427386 0.96454414 0.942156 0.75020145 0.94086496 0.94794521  
0.93236715 0.85260586 0.8783434 0.87052811]

F1 Score = [0.96705107 0.96524258 0.85965614 0.85022831 0.90147992 0.92307692  
0.94995898 0.87578419 0.85582177 0.88985633]



## 2.0.16 SVM: Fold-2

In [15]: `train_SVM(hog_img, labels, 1)`


Accuracy = 0.904

Precision = [0.98327465 0.97709321 0.82517986 0.91619203 0.9456838 0.95292988  
0.82361734 0.88880064 0.88776449 0.87326247]

Recall = [0.93005828 0.93429003 0.94950331 0.9073328 0.82001682 0.91092746  
0.99189919 0.86076923 0.83261432 0.90662139]

F1 Score = [0.9559264 0.95521236 0.88298691 0.91174089 0.87837838 0.9314554  
0.89995917 0.87456038 0.85930543 0.88962932]





## 2.0.17 SVM: Fold-3


In [20]: `train_SVM(hog_img, labels, 2)`

Accuracy = 0.9221666666666667

Precision = [0.95614789 0.93508287 0.9004894 0.90969345 0.89671737 0.92377495  
0.95446585 0.91723549 0.88569145 0.94903846]

Recall = [0.97887324 0.97973951 0.92851135 0.90444811 0.91503268 0.95140187  
0.96545616 0.87256494 0.91268534 0.81570248]

F1 Score = [0.96737712 0.95689046 0.91428571 0.9070632 0.90578245 0.9373849  
0.95992955 0.89434276 0.8989858 0.87733333]



## 2.0.18 SVM: Fold-4


In [22]: `train_SVM(hog_img, labels, 3)`

Accuracy = 0.8910833333333333

Precision = [0.99328215 0.9886823 0.84627093 0.92554194 0.90747331 0.85575365  
0.94271686 0.94048693 0.70297657 0.89053803]

Recall = [0.86178185 0.89992642 0.9197684 0.82659933 0.89395267 0.9522092  
0.9442623 0.82777778 0.94790777 0.84063047]

F1 Score = [0.92287115 0.9422188 0.88149029 0.87327701 0.90066225 0.90140845  
0.94348894 0.88054031 0.80727273 0.86486486]



## 2.0.19 SVM: Fold-5

In [24]: `train_SVM(hog_img, labels, 4)`

Accuracy = 0.8936666666666667

Precision = [0.95684647 0.95249042 0.91852487 0.87348485 0.83665644 0.87381158  
0.97885463 0.97280593 0.85244444 0.77679783]

Recall = [0.97711864 0.95984556 0.88439306 0.92461909 0.94458874 0.95557656  
0.91365132 0.63212851 0.83536585 0.91967871]

F1 Score = [0.96687631 0.95615385 0.90113589 0.89832489 0.88735258 0.91286682  
0.94512973 0.76630964 0.84381874 0.8422214 ]

